

Architecture-Dependent Failure Modes in Deep Graph Neural Networks: Separating Oversmoothing from Optimization Failure

Kiarash Ara
arak@mcmaster.ca

Yiheng Wu
wul261@mcmaster.ca

Abstract—Graph neural networks lose the ability to discriminate between nodes as depth increases, a phenomenon usually attributed to oversmoothing. Accuracy alone cannot establish that attribution: a deep model may instead fail because optimization never moves its weights. We compare GCN, GraphSAGE, GAT, and GCNII on the Cora citation network at depths 2, 4, 8, 16, and 32, recording Dirichlet energy and mean average distance (MAD) at three capture points per run, initialization, the selected checkpoint, and the final epoch, over 534 runs at ten seeds. Depth-32 failure is architecture-dependent rather than one shared mechanism. GraphSAGE and GAT do not train at depth 32: their training loss stays within 0.02 nats of the uninformed floor $\ln 7 = 1.9459$ and their checkpoint state stays near initialization. GCN partially trains, reaching 24.07% accuracy, below the 31.9% majority-class floor, and develops early-layer energy collapse with late-layer growth spanning 26 orders of magnitude across seeds. GraphSAGE’s metrics dissociate: MAD falls to near zero by depth 4 while energy plateaus near 19, because it collapses onto the constant direction, outside the augmented Laplacian’s degree-weighted null space. Jumping Knowledge restores GCN to 73.13% and GAT to 72.09% but GraphSAGE to only 40.47%; GCNII alone improves with depth, reaching 83.26%.

Index Terms—graph neural networks, oversmoothing, network depth, Dirichlet energy, node classification, trainability

I. INTRODUCTION

Many systems of practical interest are described more faithfully as graphs than as collections of independent observations: citation networks connect papers through references, recommendation systems connect users to items, and molecular graphs connect atoms through bonds. Graph neural networks (GNNs) exploit this structure through message passing, in which each node updates its representation by combining its own features with information aggregated from its neighbors [1].

Depth is the mechanism by which a GNN widens its receptive field. One message-passing layer reaches direct neighbors, two reach neighbors-of-neighbors, and an L -layer model can draw on information from roughly L hops

away. In practice, adding layers tends to make GNNs harder to train and reduces their ability to separate classes: repeated aggregation drives node representations toward one another until they become indistinguishable, a behavior commonly called oversmoothing [2, 3]. Depth is not a tunable convenience in every domain. In physics-informed graph learning, where a discretized mesh is treated as a graph, a useful prediction requires information to cross enough of the domain, so the depth at which representations degrade constrains whether the model is applicable at all.

A. The Attribution Problem

Oversmoothing is frequently invoked as though it were the only reason a deep GNN fails, but a decrease in accuracy does not establish that representations have collapsed. A deep model can also fail because optimization becomes ineffective, because gradients vanish, or because neighborhood information is compressed past usefulness. These mechanisms can produce nearly identical accuracy curves while leaving very different signatures in the hidden representations. A model that trained and then collapsed and a model whose weights never moved are separated by their internal state, not by their test accuracy.

We therefore separate two terms and use them consistently. *Oversmoothing* names a mechanism, the contraction of representations toward one another under repeated propagation. *Collapse* names what the metrics register when it happens. Oversmoothing is a claim about cause; collapse is a claim about what was measured, and establishing the first from the second requires ruling out the alternative that nothing was learned at all.

B. Approach

We study semi-supervised node classification on Cora [4, 5], holding dataset, split, and training protocol fixed so that depth is the only variable that moves. Four architectures with materially different aggregation rules are compared: GCN [6], GraphSAGE [7], GAT [8], and GCNII [9], each at depths 2, 4, 8, 16, and 32 over ten seeds.

Each run is additionally instrumented at the representation level with Dirichlet energy and mean average distance (MAD) [10], reported as a pair because they are

invariant to different things. Both are captured at three points in every run: at initialization before any gradient step, at the checkpoint selected by lowest validation loss, and at the final epoch. The epoch-0 capture isolates the architecture from the optimizer, and the epoch-0-to-checkpoint comparison is what separates a trained state from an untrained one.

C. Contributions

- **An architecture-resolved account of depth-32 failure**, showing that GCN, GraphSAGE, and GAT fail through distinct mechanisms rather than one shared collapse, separated by measurements that do not depend on accuracy.
- **A three-capture diagnostic protocol** distinguishing collapse present at initialization from collapse that develops during training, applied uniformly across 534 runs.
- **A measured explanation of metric dissociation**: GraphSAGE collapses onto the constant direction, which lies outside the augmented Laplacian’s degree-weighted null space, so its energy cannot vanish however completely its directions align.
- **An ablation of three mitigations and their combination** at depth 32, showing that the most effective one is architecture-dependent rather than universal.

The study is confined to Cora and its standard public split, so the depth at which performance declines and the ranking of mitigations should not be assumed to transfer. Per-layer gradient norms are not recorded, so explanations appealing to gradient behavior are stated as consistent with the measurements rather than established by them.

II. RELATED WORK

A. Message Passing and Depth

Most modern GNNs are instances of the message-passing framework [1], differing in the aggregation rule. GCN applies a symmetrically normalized adjacency operator [6]; GraphSAGE transforms the node and the neighborhood mean through separate weight matrices and was introduced for inductive settings [7]; GAT replaces fixed degree weights with learned attention [8]. Because depth determines how far information travels, these differing rules are also differing answers to what repeated application does to a representation.

B. Oversmoothing and Its Measurement

Li, Han, and Wu identified graph convolution as a form of Laplacian smoothing and showed that the operation letting neighbors share information eventually makes them indistinguishable [2]. Oono and Suzuki strengthened this into an asymptotic statement: under conditions on the weight spectra, deep GNN representations approach a low-dimensional invariant subspace exponentially in depth [3]. Cai and Wang recast the phenomenon in terms of Dirichlet

energy, which contracts multiplicatively at a rate governed by the weight norms and the graph spectral gap [11]. These results establish that contraction happens; they do not establish that a given model’s poor accuracy was caused by it.

Chen et al. propose MAD, the mean pairwise cosine distance among node representations, together with MADGap, which correlates with test accuracy closely enough to serve as a training-free quality predictor [10]. That correlation is precisely why this study reports MAD but not MADGap: a metric validated as an accuracy predictor cannot then be used to ask whether collapse and poor accuracy are separable.

C. Trainability

Depth degrades GNNs through routes other than smoothing. Kipf and Welling report a depth study in an appendix rather than in their main results, finding best performance at two to three layers and training becoming difficult beyond roughly seven layers without residual connections [6]. That is a statement about optimization, not about representation collapse, and it is the prior finding this work extends.

D. Mitigations

Residual connections add an identity path that preserves an earlier representation and shortens the gradient route. PairNorm instead acts on the geometry directly, centering and rescaling rows so their dispersion does not shrink with depth [12]. Jumping Knowledge changes the readout rather than the propagation, combining representations from several depths [13]. GCNII modifies the layer itself, adding an initial-residual term and an identity mapping [9].

E. The Gap

Prior work establishes that accuracy declines with depth and, separately, that representations contract with depth. What is rarely established is that the first was caused by the second in a particular model. Evaluations typically report a single collapse metric at a single point in training, usually the trained model, which cannot separate a network that trained and then collapsed from one whose weights never left initialization. This study addresses that gap by holding dataset, split, and protocol fixed across four architectures, reporting two metrics with different invariances as a pair, and capturing both at three points so the untrained state is an explicit reference.

III. METHODOLOGY

A. Data and Task

All experiments use the Cora citation network through PyTorch Geometric’s `Planetoid` loader [4, 5]: 2,708 papers, 1,433-dimensional binary bag-of-words features, seven classes, and the standard public split of 140 training, 500 validation, and 1,000 test nodes. The training set

holds 20 labeled papers per class. Features are row-normalized, matching the published GCN setup [6]. Every node participates in message passing regardless of split; only training-node labels enter the loss.

Let $G = (V, E)$ with $|V| = N$, feature matrix $X \in \mathbb{R}^{N \times F}$, and adjacency A . Self-loops are added, $\tilde{A} = A + I$, with augmented degrees $\tilde{d}_i = \sum_j \tilde{A}_{ij}$ and $\tilde{D} = \text{diag}(\tilde{d}_1, \dots, \tilde{d}_N)$. The normalized propagation operator and Laplacian are

$$\hat{A} = \tilde{D}^{-1/2} \tilde{A} \tilde{D}^{-1/2}, \quad \mathcal{L} = I - \hat{A}.$$

B. Architectures

Four update rules are compared. GCN [6] applies

$$H^{(\ell+1)} = \sigma(\hat{A}H^{(\ell)}W^{(\ell)}), \quad H^{(0)} = X.$$

GraphSAGE with mean aggregation [7] transforms the node and its neighborhood mean $m_i^{(\ell)} = |\mathcal{N}(i)|^{-1} \sum_{j \in \mathcal{N}(i)} h_j^{(\ell)}$ separately,

$$h_i^{(\ell+1)} = \sigma(W_{\text{self}}^{(\ell)} h_i^{(\ell)} + W_{\text{neigh}}^{(\ell)} m_i^{(\ell)} + b^{(\ell)}).$$

GAT [8] replaces fixed weights with attention. With $q_i^{(\ell)} = W^{(\ell)} h_i^{(\ell)}$ and $e_{ij}^{(\ell)} = \text{LeakyReLU}((a^{(\ell)})^T [q_i^{(\ell)} \| q_j^{(\ell)}])$,

$$\alpha_{ij}^{(\ell)} = \frac{\exp(e_{ij}^{(\ell)})}{\sum_{k \in \mathcal{N}(i) \cup \{i\}} \exp(e_{ik}^{(\ell)})}.$$

GCNII [9] adds an initial residual and an identity mapping,

$$H^{(\ell+1)} = \sigma\left(\left[(1 - \alpha_\ell)\hat{A}H^{(\ell)} + \alpha_\ell H^{(0)}\right]\left[(1 - \beta_\ell)I + \beta_\ell W^{(\ell)}\right]\right)$$

with $\alpha = 0.1$ and $\theta = 0.5$ from the original paper. Its two required linear projections are placed outside the counted depth L , so that L denotes the same number of message-passing hops for all four architectures. The known cost is two extra parameterized layers, bounded in size and independent of depth, which can only overstate GCNII's capacity relative to its peers.

All four are realized as one parameterized model in which `convType` selects the convolution and depth, width, mitigation hooks, and the recording position are otherwise identical, so a measured difference is attributable to the convolution. Hidden width is 64 throughout rather than each architecture's published width, because capacity parity is a precondition for comparing at matched depth: at width 16, GAT's eight heads would attend over two-dimensional keys.

C. Mitigations

Three mitigations attach to the shared layer loop by composition. A residual connection adds an identity path, $H^{(\ell+1)} = F^{(\ell)}(H^{(\ell)}) + H^{(\ell)}$, applicable only where widths match. PairNorm [12] centers and rescales rows,

$$\tilde{H} = H - \frac{1}{N} \mathbf{1}\mathbf{1}^T H, \quad h_i^{\text{PN}} = s \frac{\tilde{h}_i}{\|\tilde{h}_i\|_2 + \varepsilon},$$

the scale-individual variant (PN-SI). Jumping Knowledge (JK) [13] changes the readout, pooling across depths elementwise,

$$H_{\text{JK}} = \max(H^{(1)}, H^{(2)}, \dots, H^{(L)}),$$

followed by a linear classifier.

D. Representation Diagnostics

For an embedding matrix H with rows h_i , the normalized-Laplacian Dirichlet energy is

$$E(H) = \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \tilde{A}_{ij} \left\| \frac{h_i}{\sqrt{\tilde{d}_i}} - \frac{h_j}{\sqrt{\tilde{d}_j}} \right\|_2^2 = \text{Tr}(H^T \mathcal{L} H),$$

reported per dimension as $E_{\text{dim}}(H) = E(H)/\text{dim}(H)$ so that layers of differing width remain comparable. MAD is the mean cosine distance over node pairs with nonzero distance [10].

The two are reported together because they are invariant to different things. Cosine distance is invariant to positive per-node rescaling: for $S = \text{diag}(s_1, \dots, s_N)$ with $s_i > 0$, $\text{MAD}(SH) = \text{MAD}(H)$, so MAD captures directional alignment and discards magnitude. Dirichlet energy is quadratic, $E_{\text{dim}}(cH) = c^2 E_{\text{dim}}(H)$, and retains both. Their disagreement is therefore informative: representations can align almost completely, driving MAD toward zero, while energy stays high because the direction they align on is not the Laplacian's degree-weighted null direction.

Not every stored layer is comparable. $H^{(0)}$ is excluded as sparse 1,433-dimensional input, and under a last-layer readout $H^{(L)}$ is excluded as seven-dimensional logits that cross-entropy shapes directly. The comparable band is therefore $\mathcal{B} = \{1, \dots, L-1\}$, or $\{1, \dots, L\}$ under Jumping Knowledge. Energy is reported relative to the band's first layer, $R_\ell = E_{\text{dim}}(H^{(\ell)})/E_{\text{dim}}(H^{(\ell_0)})$ with $\ell_0 = \min(\mathcal{B})$, and summarized by a contraction slope b fitted as

$$\log \max\{E_{\text{dim}}(H^{(\ell)}), \varepsilon_E\} \approx a + b\ell, \quad \varepsilon_E = 10^{-12},$$

the floor guarding against $\log 0$. The floor can pull a slope toward zero but cannot create a trend that is not present.

E. Three Capture Points

Every run is measured three times, each an eval-mode forward pass under `torch.no_grad()`: at epoch 0 before any gradient step, at the final epoch on the weights the loop ended with, and at the selected checkpoint after the best validation-loss state is restored. The final-epoch capture is taken *before* the restore; taken afterward it would duplicate the checkpoint capture on every run. The epoch-0 capture

isolates the architecture from the optimizer, and the epoch-0-to-checkpoint comparison is what distinguishes a model that trained and collapsed from one that never trained.

F. Training Protocol and Experiment Grid

All runs use Adam [14] with learning rate 0.01, dropout 0.5, and weight decay 5×10^{-4} , the configuration selected once by a depth-2 GCN search and then frozen across every architecture and depth, so that the depth effect is not confounded with a hyperparameter effect. Weight decay is applied uniformly across layers rather than to the first layer only, so regularization strength does not become an implicit function of depth. Early stopping and checkpoint selection are both driven by validation loss, with patience 100 and a ceiling of 1,000 epochs. Patience is held constant across depth: at a short window, a 32-layer run whose loss plateaus early halts near epoch 12, leaving a flat curve that a slow optimization and an untrainable architecture would share.

The sweep is organized as six arms totaling 534 runs: an unmitigated sweep over GCN, GraphSAGE, and GAT (150 runs); a mitigation ablation on GCN over residual, PairNorm, JK, and PairNorm+residual (200); GCNII (50); the winning mitigation carried to GraphSAGE and GAT (100); a fidelity arm at the published width 16 (10); and the hyperparameter search (24). Depths are $\{2, 4, 8, 16, 32\}$, log-spaced because the contraction bound predicts exponential decay in depth, and every configuration is repeated over ten seeds controlling both initialization and training stochasticity. All runs are CPU, Python 3.14.4, PyTorch 2.13.0, PyTorch Geometric 2.8.0.

IV. EXPERIMENTS AND RESULTS

A. Baseline Reproduction

Every reported number is regenerated from the stored per-run records, and results are mean and standard deviation over ten seeds. Before depth is varied, each architecture is checked at depth 2 against the figure its own paper reports. GCN reaches 81.49% (± 0.32) against 81.5% [6]; GraphSAGE reaches 80.10% (± 0.42), for which no published Cora figure exists [7]; GAT reaches 78.90% (± 1.57) against 83.0% (± 0.7) [8], a shortfall consistent with the shared configuration, from which its own published setup differs most.

B. Accuracy Across Depth

All three conventional architectures degrade sharply with depth (Fig. 1). The drop is not gradual: each loses 4.4 to 8.4 points between depths 2 and 4, falls steeply between 4 and 8, and stays low. At depth 32, GCN reaches 24.07%, GraphSAGE 22.65%, and GAT 19.66%, all below the 31.9% majority-class floor computed from the test-split label counts; GCN’s full ten-seed range there, 21.2% to 26.2%, lies entirely below it.

Macro-F1 collapses faster still: at depth 32 GAT’s is 0.0459 and GraphSAGE’s 0.0514, against GCN’s 0.2102,




Fig. 1: Test accuracy versus depth, unmitigated GCN, GraphSAGE, and GAT.

consistent with predictions concentrating onto a few classes rather than spreading evenly across all seven.

C. Collapse at Initialization

At epoch 0 the weights are random, so any contraction is attributable to the architecture rather than to training. GCN and GAT collapse cleanly in both direction and magnitude, while GraphSAGE collapses in direction only (Table I). GCN’s energy ratio falls about 12 orders of magnitude by depth 32 and GAT’s about 10, consistent with the multiplicative contraction Cai and Wang describe [11], though its numerical tightness is not tested here.

GraphSAGE’s two metrics dissociate at every depth: MAD reaches near zero by depth 4 while its energy ratio *rises* and then plateaus near 19. A singular-value decomposition at depth 16 over three seeds locates the cause. GraphSAGE’s representation matrix is exactly rank-1 (top-singular-value energy fraction 1.0000), and its top singular vector matches the constant direction at cosine 1.0000 against 0.9512 to the $\sqrt{\text{degree}}$ direction; GCN and GAT are only mostly rank-1 (0.93 to 0.98) and do not separate the two candidates. Because Cora’s degrees run from 2 to 169, the augmented Laplacian’s null space is the $\sqrt{\text{degree}}$ direction rather than the constant vector, so GraphSAGE collapses onto a direction outside that null space and its energy cannot vanish however completely its representations align.

D. Separating Trained from Untrained Failure

Training changes the picture, and differently per architecture. At the selected checkpoint GCN’s MAD no longer falls with depth, holding near 0.183 at depth 32, while GraphSAGE’s and GAT’s remain essentially zero, close to their own epoch-0 values.

Three independently measured quantities converge for GraphSAGE and GAT. Their training loss never leaves the

TABLE I: Epoch-0 MAD and relative energy versus depth, means over ten seeds. At depth 2 the comparable band holds one layer, so the ratio is not reported.

Conv	Metric	$d=2$	$d=4$	$d=8$	$d=16$	$d=32$
GCN	MAD	0.600	0.240	0.085	0.023	0.0045
GCN	E_{last}/E_1	—	3.1e-2	3.1e-4	2.6e-7	6.2e-13
GAT	MAD	0.596	0.211	0.076	0.017	0.0041
GAT	E_{last}/E_1	—	7.3e-2	2.7e-3	1.3e-5	1.3e-10
SAGE	MAD	0.083	1.8e-4	9.9e-7	≈ 0	≈ 0
SAGE	E_{last}/E_1	—	18.3	19.8	20.3	19.3

uninformed floor $\ln 7 = 1.9459$, reaching minima of 1.9343 and 1.9393, a movement of 0.01 to 0.02 nats (Fig. 2); their runs exhaust the patience of 100 epochs almost immediately, at 102 to 168 and 101 to 121 epochs; and their checkpoint MAD is essentially identical to their epoch-0 MAD. None is derived from the others, which makes their agreement convergence rather than circularity: at depth 32 these two architectures do not train at all.




Fig. 2: Training loss versus epoch at depth 32, unmitigated GCN against GCN with Jumping Knowledge; the dotted line marks $\ln 7$.

GCN behaves differently. Its loss descends to a mean minimum of 1.57 and its runs last 315 to 883 epochs. Its checkpoint energy ratio grows rather than decays, spanning 26 orders of magnitude across seeds at depth 32, because the *denominator* has collapsed: E_1 falls to approximately $1.4e-15$ and as low as $2.9e-38$. For nine of ten seeds the first 22 to 28 of 31 band layers carry energy below the 10^{-12} floor while the last few carry very large energy. Depth-32 failure therefore separates into two training modes, with distinct collapse signatures inside the non-training pair.

E. Mitigations

Jumping Knowledge is the clear winner on GCN, and the ranking does not carry across architectures (Table II). Residual alone does not beat the unmitigated baseline at depth 32, and PairNorm combined with residual is worse than PairNorm alone. Carried to the other architectures,

TABLE II: Depth-32 test accuracy for every arm, ten seeds each. GCNII reaches its best result at the greatest depth tested.

Configuration	Depth-32 acc.	Macro-F1
GCN, unmitigated	24.07% (± 1.60)	0.2102
GCN + residual	19.25% (± 8.97)	—
GCN + PairNorm	58.07% (± 7.10)	—
GCN + PairNorm + residual	24.36% (± 9.79)	—
GCN + JK	73.13% (± 3.43)	0.7249
GraphSAGE, unmitigated	22.65% (± 9.83)	0.0514
GraphSAGE + JK	40.47% (± 3.76)	0.2611
GAT, unmitigated	19.66% (± 8.46)	0.0459
GAT + JK	72.09% (± 3.77)	0.7252
GCNII	83.26% (± 0.44)	—

JK recovers GAT almost to GCN’s level including macro-F1 (0.7252 against 0.7249), but recovers GraphSAGE only to 40.47% with macro-F1 0.2611. All ten GraphSAGE+JK seeds sit above the majority-class floor, so the gain is real, but it is about a third of what the same mitigation achieves elsewhere and is not spread evenly across classes. GCNII is the only architecture whose accuracy improves with depth, and its checkpoint contraction slope stays within roughly ± 0.4 at every depth (+0.383, +0.169, -0.037 , +0.012) where GCN’s runs from +1.64 to +2.46.

V. CONCLUSION AND FUTURE WORK

Depth-related failure in graph neural networks is architecture-dependent, and accuracy alone cannot show this. Measuring Dirichlet energy and MAD at three capture points across 534 runs separates three behaviors that a single accuracy curve would render identical. GraphSAGE and GAT do not train at depth 32: their loss never leaves $\ln 7$, patience exhausts almost immediately, and their trained state is close to their untrained one, so attributing their failure to oversmoothing would be incorrect. GCN partially trains and develops a distinct signature of early-layer energy collapse paired with late-layer growth. GCNII alone improves with depth, reaching 83.26% at depth 32.

Two results depend on the paired-metric design. GraphSAGE’s MAD and energy disagree at every depth because it collapses onto the constant direction, which lies outside the degree-weighted null space of Cora’s augmented Laplacian; either metric alone would have misled. And Jumping Knowledge, the best mitigation on GCN, transfers to GAT but not to GraphSAGE, so the countermeasure cannot be chosen independently of the architecture.

The study is limited to one dataset and one split, and per-layer gradient norms are not recorded, so the mechanism inferred for GCN’s early-layer collapse remains consistent with the measurements rather than established by them. Adding that instrumentation is the most direct next step. Two experiments would isolate why GraphSAGE collapses onto the constant direction at one hop: a degree-normalized-mean variant and a Glorot-initialized [15] SAGEConv, run separately. Extending the protocol to graphs with other degree distributions would show whether the null-space result is Cora-specific. The protocol, rather than the numbers, is what is expected to transfer.

REFERENCES

- [1] J. Gilmer, S. S. Schoenholz, P. F. Riley, O. Vinyals, and G. E. Dahl, “Neural message passing for quantum chemistry,” in *Proceedings of the 34th International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, vol. 70. PMLR, 2017, pp. 1263–1272. [Online]. Available: <https://proceedings.mlr.press/v70/gilmer17a.html>
- [2] Q. Li, Z. Han, and X.-M. Wu, “Deeper insights into Graph Convolutional Networks for semi-supervised learning,” *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 32, no. 1, pp. 3538–3545, 2018. [Online]. Available: <https://doi.org/10.1609/aaai.v32i1.11604>
- [3] K. Oono and T. Suzuki, “Graph Neural Networks exponentially lose expressive power for node classification,” in *8th International Conference on Learning Representations*. OpenReview.net, 2020. [Online]. Available: <https://openreview.net/forum?id=S1ldO2EFP>
- [4] P. Sen, G. Namata, M. Bilgic, L. Getoor, B. Gallagher, and T. Eliassi-Rad, “Collective classification in network data,” *AI Magazine*, vol. 29, no. 3, pp. 93–106, 2008. [Online]. Available: <https://doi.org/10.1609/aimag.v29i3.2157>
- [5] Z. Yang, W. W. Cohen, and R. Salakhutdinov, “Revisiting semi-supervised learning with graph embeddings,” in *Proceedings of the 33rd International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, vol. 48. PMLR, 2016, pp. 40–48. [Online]. Available: <https://proceedings.mlr.press/v48/yanga16.html>
- [6] T. N. Kipf and M. Welling, “Semi-supervised classification with Graph Convolutional Networks,” in *5th International Conference on Learning Representations*. OpenReview.net, 2017. [Online]. Available: <https://openreview.net/forum?id=SJU4ayYgl>
- [7] W. L. Hamilton, Z. Ying, and J. Leskovec, “Inductive representation learning on large graphs,” in *Advances in Neural Information Processing Systems*, vol. 30. Curran Associates, Inc., 2017, pp. 1024–1034. [Online]. Available: https://papers.nips.cc/paper_files/paper/2017/hash/5dd9db5e033da9c6fb5ba83c7a7e9bea9-Abstract.html
- [8] P. Veličković, G. Cucurull, A. Casanova, A. Romero, P. Liò, and Y. Bengio, “Graph Attention Networks,” in *6th International Conference on Learning Representations*. OpenReview.net, 2018. [Online]. Available: <https://openreview.net/forum?id=rJXmpikCZ>
- [9] M. Chen, Z. Wei, Z. Huang, B. Ding, and Y. Li, “Simple and deep Graph Convolutional Networks,” in *Proceedings of the 37th International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, vol. 119. PMLR, 2020, pp. 1725–1735. [Online]. Available: <https://proceedings.mlr.press/v119/chen20v.html>
- [10] D. Chen, Y. Lin, W. Li, P. Li, J. Zhou, and X. Sun, “Measuring and relieving the over-smoothing problem for Graph Neural Networks from the topological view,” *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 34, no. 4, pp. 3438–3445, 2020. [Online]. Available: <https://doi.org/10.1609/aaai.v34i04.5747>
- [11] C. Cai and Y. Wang, “A note on over-smoothing for Graph Neural Networks,” *arXiv preprint arXiv:2006.13318*, 2020. [Online]. Available: <https://arxiv.org/abs/2006.13318>
- [12] L. Zhao and L. Akoglu, “PairNorm: Tackling oversmoothing in GNNs,” in *8th International Conference on Learning Representations*. OpenReview.net, 2020. [Online]. Available: <https://openreview.net/forum?id=rkecl1rtwB>
- [13] K. Xu, C. Li, Y. Tian, T. Sonobe, K.-i. Kawarabayashi, and S. Jegelka, “Representation learning on graphs with Jumping Knowledge Networks,” in *Proceedings of the 35th International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, vol. 80. PMLR, 2018, pp. 5453–5462. [Online]. Available: <https://proceedings.mlr.press/v80/xu18c.html>
- [14] D. P. Kingma and J. Ba, “Adam: A method for stochastic optimization,” in *3rd International Conference on Learning Representations*, 2015. [Online]. Available: <https://arxiv.org/abs/1412.6980>
- [15] X. Glorot and Y. Bengio, “Understanding the difficulty of training deep feedforward neural networks,” in *Proceedings of the Thirteenth International Conference on Artificial Intelligence and Statistics*, ser. Proceedings of Machine Learning Research, vol. 9, 2010, pp. 249–256. [Online]. Available: <https://proceedings.mlr.press/v9/glorot10a.html>